

# Lecture 21: Priority Queues & Heaps

PIC 10B  
Todd Wittman




## Priority Queues

- Suppose we are keeping track of a line of patients entering a hospital emergency room.
- In a queue, the first person to enter the emergency room is the first person to see the doctor (FIFO).
- But what if a patient comes in who needs immediate medical attention (priority 1)?
- In a priority queue, each element is assigned a unique priority.
- Elements with the highest priority should be the first to leave the queue (HPFO?).  
3, 2, 8 1, 5, 6, 7 The first one to leave is the 1.
- There are several ways to implement a priority queue. We will use the data structure called a heap.



## The Heap Data Structure


- A heap (or *min-heap*) is a data structure that is a cross between a queue and a binary search tree. Essentially it has the shape of a binary tree and the ordering of a queue.
- Each node of the heap stores a value and has two children, just like in a binary tree.



- **The Heap Property:** Every node has a value smaller than both its children.
- **The Shape Property:** Every level of the heap is full, except for possibly the last level which is filled left to right.
- **The Queue Property:** Elements can be removed or accessed only at the root node. (No traversals.)

## The Heap Concept


- Here's an example of 2 possible heaps on the numbers 0-8.



- Note that the Heap Property implies that the minimum value is always at the root.
- So the min will be the first value to leave the heap.
- The Shape Property implies that tree is always perfectly balanced.
- Every heap with same number of nodes has the same shape.

# Heaps vs. Binary Search Trees


- They may look similar, but the ordering is very different.



- The BST may contain some blank spots.
- For a BST, the height  $h$  may vary:  $\lfloor \log_2 N \rfloor \leq h \leq N - 1$
- For a heap, the height  $h$  is always:  $h = \lfloor \log_2 N \rfloor$
- You can traverse all the nodes of a tree. In a heap, you can only look at the root.

## Entering a Heap

- To enter a heap, we put the new value in the next position at the bottom of the heap.
- To maintain the Heap Property, we then "bubble" or "sift" the value up the heap by swapping the value with its parent if necessary.




- Note we have to swap at most  $h$  nodes, so inserting is  $O(\log N)$ .

## Building A Heap

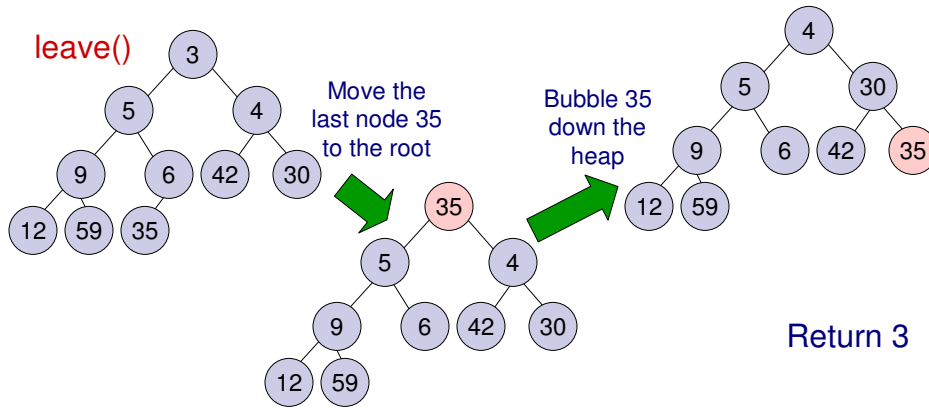
- To create a heap, we repeatedly insert the new values.
- Ex Build a heap by inserting the following values in order:

**4, 8, 3, 1, 5, 2**




## Leaving a Heap

- Only the root node can be removed from a heap.
- We swap the root value with the value in the last position of the heap and delete the last node.
- We then bubble this value down the tree, swapping with the smaller of its 2 children. Note this is also a  $O(\log N)$  operation.



## Implementing a Heap with a Vector

- Because the heap is full, we can implement a heap using a vector rather than using a tree and pointers.
- Read the heap in a level-order traversal, reading each row left to right. Note the root is at  $v[0]$ .



- For the element at  $v[i]$  we have:
  - Left child:  $v[2i+1]$
  - Right child:  $v[2i+2]$
  - Parent:  $v[(i-1)/2]$



## The Heap Class

- The Heap class declaration looks a lot like the Queue class.

```
template <typename T>
class Heap {
public:
    Heap();
    bool isEmpty() const;
    T peek() const;
    void enter(T value);
    T leave();
private:
    vector<T> v;
    void swap(int i, int j);
};
```

- Do we need the Big 4?



## Basic Heap Functions

```
template <typename T>
Heap<T>::Heap() {
    v.resize(0);
}
```

```
template <typename T>
T Heap<T>::peek() const {
    return v[0];
}
```

```
template <typename T>
bool Heap<T>::isEmpty() const {
    return v.size()==0;
}
```

```
template <typename T>
void Heap<T>::swap(int i, int j) {
    T temp = v[i];
    v[i] = v[j];
    v[j] = temp;
    return;
}
```

## Entering a Heap

```
template <typename T>
void Heap<T>::enter(T value) {
    v.push_back(value);
    int pos = v.size()-1;
    int parent = (int) (pos-1)/2;
    while (parent >= 0 && v[pos] < v[parent]) {
        swap(parent, pos);
        pos = parent;
        parent = (int) (pos-1)/2;
    }
    return;
}
```

Add new value to the bottom of the heap.

Bubble the new value up by swapping with its parent.

## Leaving a Heap

```
template <typename T>
```

```
T Heap<T>::leave() {
```

```
    T top = v[0];
```

Store the value at the root. We return this at the end.

```
    v[0] = v[v.size()-1];
```

Copy value at bottom into root and delete last node.

```
    v.pop_back();
```

```
    int pos = 0;
```

```
    bool continueBubbleDown = (2*pos+1 < v.size());
```

```
    while (continueBubbleDown) {
```

```
        if (2*pos+1 >= v.size())
```

Case 1: No children means stop.

```
            continueBubbleDown = false;
```

```
        else if (2*pos+2 >= v.size()) {
```

```
            if (v[pos] > v[2*pos+1]) {
```

```
                swap(pos, 2*pos+1);
```

```
                pos = 2*pos+1;
```

Case 2: Only left child.  
We have reached bottom of heap, so perform at most 1 more swap.

```
            }
```

```
            continueBubbleDown = false;
```

```
        }
```

More....

## Leaving a Heap

```
    else
```

Case 3: 2 children

```
        if (v[pos] > v[2*pos+1] && v[2*pos+1] <= v[2*pos+2]) {
```

```
            swap(pos, 2*pos+1);
```

```
            pos = 2*pos+1;
```

Left child is smaller.

```
        }
```

```
        else if (v[pos] > v[2*pos+2]) {
```

```
            swap(pos, 2*pos+2);
```

```
            pos = 2*pos+2;
```

Right child is smaller.

```
        }
```

```
        else
```

If both children are bigger, stop.

```
            continueBubbleDown = false;
```

```
    }
```

```
    return top;
```

Return the value that was at the root.

```
}
```

## Example Program

- Put 25 random 4-letter words into a heap and then remove them one at a time.

```
int main() {  
    Heap<string> H;  
    string word = "aaaa";  
    for (int i=0; i<25; i++) {  
        for (int j=0; j<4; j++)  
            word[j] = (char) (rand()%26+(int)'a');  
        H.enter(word);  
    }  
    while (!H.isEmpty())  
        cout << H.leave() << "\n";  
    return 0;  
}
```



- Note this prints the words in sorted order from smallest to largest.

## Heap Sort

- We can sort the elements in a vector by placing the values into a heap and removing them one at a time.
- Each enter/leave is  $O(\log N)$ , so doing  $N$  inserts/leaves takes  $O(N \log N)$  time.
- So we can heap sort a list in  $O(N \log N)$  time.
- But this requires an additional  $O(N)$  memory for the heap vector.
- There is a special way to "heapify" a given vector, so that the sorting is done *in place*.
- You will implement heap sort using this trick for this week's homework. Do not use the Heap class for heap sort.